



Deep Learning Meets Data Privacy

Presenter: Sharika Loganathan
Data Scientist,
University of Pittsburgh

Agenda

Why Privacy Matters in the Age of AI

Limitations of Traditional Anonymization

From Anonymization to Computable Privacy

Understanding Differential Privacy

Rethinking the Machine Learning Pipeline

Differential Privacy in Practice

Q&A, Feedback, and References

Why Privacy Matters in the Age of Big Data & AI



Massive Data
Collection



AI Models
Learn from Us



Risk of Re-
identification



Trust &
Transparency



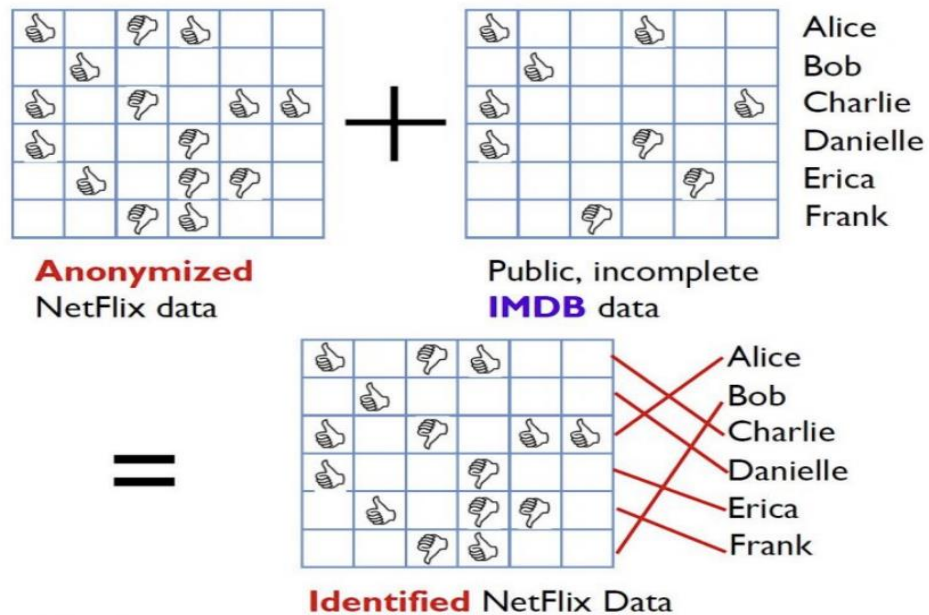
Legal & Ethical
Imperatives



Sustainable AI
Development

According to Statista, the global data volume was 149 zettabytes in 2024 and is projected to reach 181 zettabytes by the end of 2025

The Problem: Why Traditional Anonymization Fails



A dataset with **Date of Birth, Sex, and ZIP code** can re-identify most U.S. citizens.

In fact, **87%** of Americans can be uniquely identified using just these three attributes

Linkage Attacks - Netflix's anonymized movie ratings were de-anonymized using IMDb reviews.

Traditional anonymization techniques fail against modern data linkage and AI-driven inference attacks — demanding stronger privacy guarantees like **Differential Privacy**.

Need for Computational Privacy

GOAL : Privacy Preserving Data Analysis

- Obtain answers or usefulness in surveys, but at the same time maintain the privacy or “plausible deniability
- Provides **privacy guarantees**
- Protects against **a wide range of attacks** — even unforeseen ones
- Ensures **trustworthy data analysis** in deep learning and AI



What is Differential Privacy



Definition:

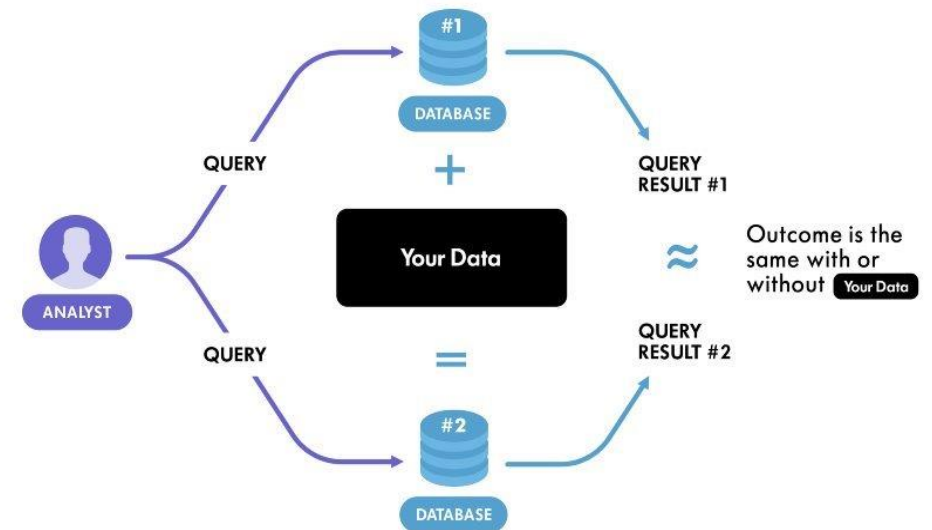
Differential privacy (DP) is a mathematically rigorous framework for releasing statistical information about datasets while protecting the privacy of individual data subjects.

Key Idea (Cynthia Dwork, 2016):

“The outcome of any analysis is essentially equally likely, independent of whether any individual joins or refrains from joining the dataset.”

Core Principle (Kearns & Roth, 2020):

"Adding or removing the data record of a single individual does not change the probability of any result too much".



Google receives only noisy data and aggregates it...

Browser	True Count	Private Count	Error	
Chrome	6,483	6,956	7.3	%
Safari	2,055	2,799	36.2	%
Edge	753	1,622	115.4	%
Firefox	505	1,590	214.9	%
Other	204	1,202	489.2	%

=====

KEY INSIGHTS

=====

- ✔ Individual Privacy: No one (not even Google) knows your actual browser
- ✔ Useful Statistics: Aggregate trends are preserved with high accuracy
- ✔ Mathematical Guarantee: ϵ -differential privacy ensures plausible deniability

- 💡 With $\epsilon=1.0$:
- Any individual can deny their true browser choice
 - Aggregate statistics have ~1.0% statistical error

=====

PRIVACY GUARANTEE DEMONSTRATION

=====

Example: Alice uses Firefox, but reports 'Chrome' due to randomization

Alice's TRUE browser: firefox
Alice's 10 reports: ['safari', 'firefox', 'chrome', 'safari', 'other', 'chrome', 'edge', 'chrome', 'firefox', 'safari']

→ Even if someone intercepts Alice's report, they can't be sure of her actual browser because randomization provides plausible deniability!

=====

BONUS: RAPPOR – GOOGLE'S ACTUAL CHROME IMPLEMENTATION

=====

Browser: chrome

5 RAPPOR-encoded reports (each is a bit vector):

Report 1: [1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0]

Report 2: [0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1]

Report 3: [1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1]

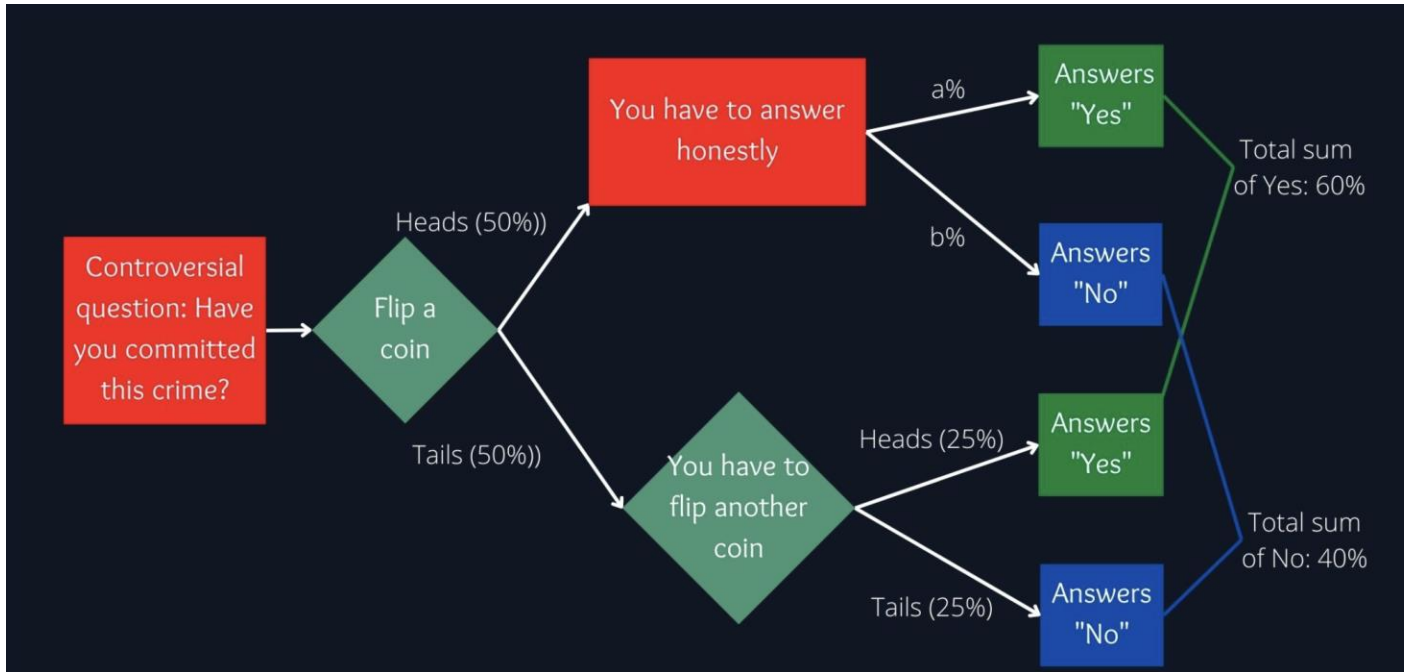
Report 4: [0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1]

Report 5: [0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1]


Note: RAPPOR provides even stronger privacy by:

- Using Bloom filters for efficient encoding
- Adding multiple layers of randomization
- Enabling longitudinal privacy (multiple reports from same user)


Mathematics Behind Differential Privacy



Laplace Noise



Gaussian Noise



$$\frac{\text{Probability of seeing output } O \text{ on input } D_1}{\text{Probability of seeing output } O \text{ on input } D_2} = \frac{\Pr[\mathcal{M}(D_1) \in O]}{\Pr[\mathcal{M}(D_2) \in O]} \leq e^\epsilon$$

Indistinguishability: bounded ratio of probabilities

Mathematical Formula:

$$TY = HY + FY$$

where:

TY = Total Yes (observed)

HY = Honest Yes = $p \times \text{True_Yes_Rate}$

FY = False Yes = $(1-p) \times 0.5 = 0.25$

Solving for True_Yes_Rate:

$$\begin{aligned}\text{True_Yes_Rate} &= (TY - FY) / p \\ &= (0.6010 - 0.25) / 0.5 \\ &= 0.7020\end{aligned}$$

Metric	Value
True 'Yes' Rate	70.00%
Observed 'Yes' Rate	60.10%
Estimated 'Yes' Rate	70.20%
Estimation Error	0.20%

Press ENTER to verify privacy guarantees...

=====

DIFFERENTIAL PRIVACY VERIFICATION

=====

 Testing the mathematical guarantee:


$$\Pr[K(D_1) \in S] \leq \exp(\epsilon) \times \Pr[K(D_2) \in S] + \delta$$

Dataset D₁: 600 'Yes' → Pr = 0.6000


Dataset D₂: 601 'Yes' → Pr = 0.6010

Privacy Loss: $\ln(0.6010/0.6000) = 0.0017$

Required Bound: $\epsilon + \delta = 1.0000 + 1e-05 = 1.0000$

 SATISFIES (ε,δ)-Differential Privacy

Privacy Budget Epsilon

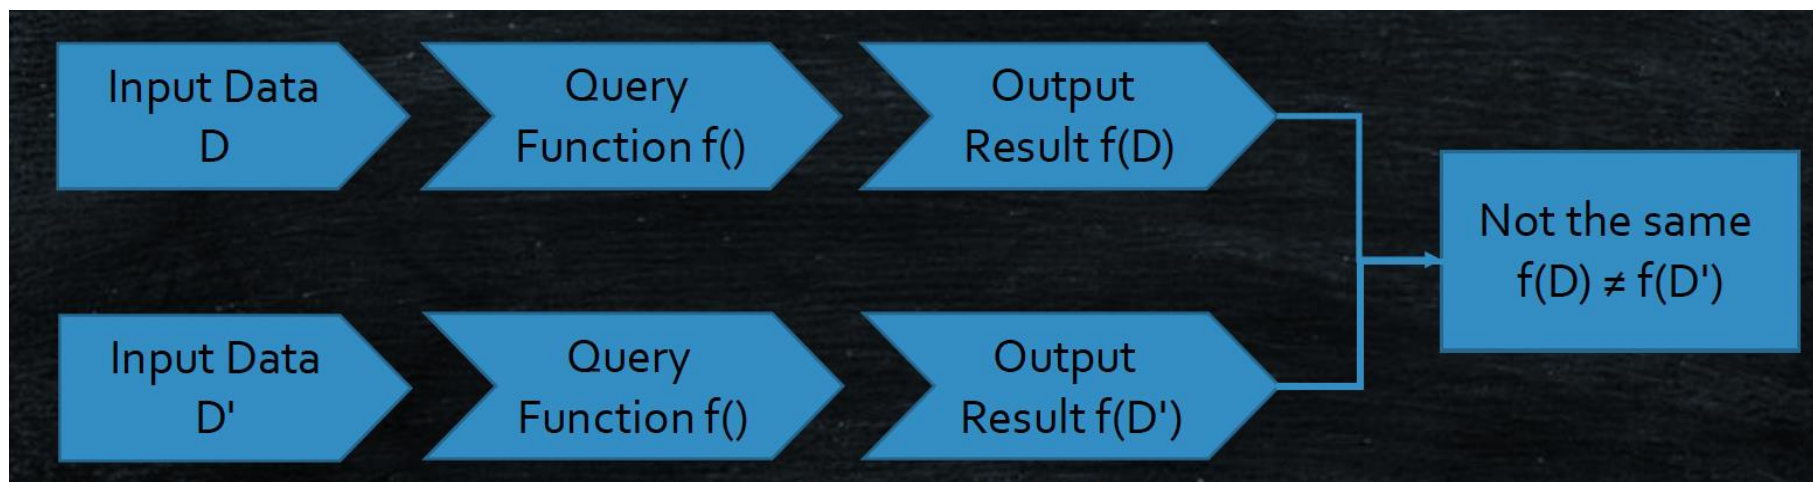
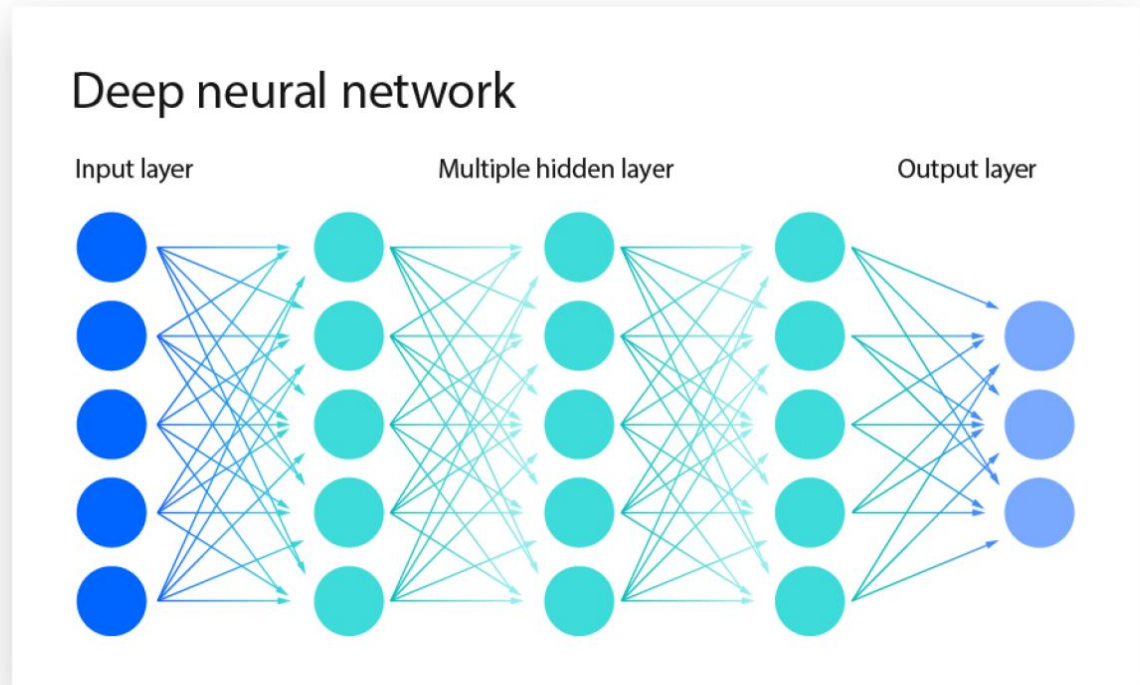


Initial suspicion was 50%,

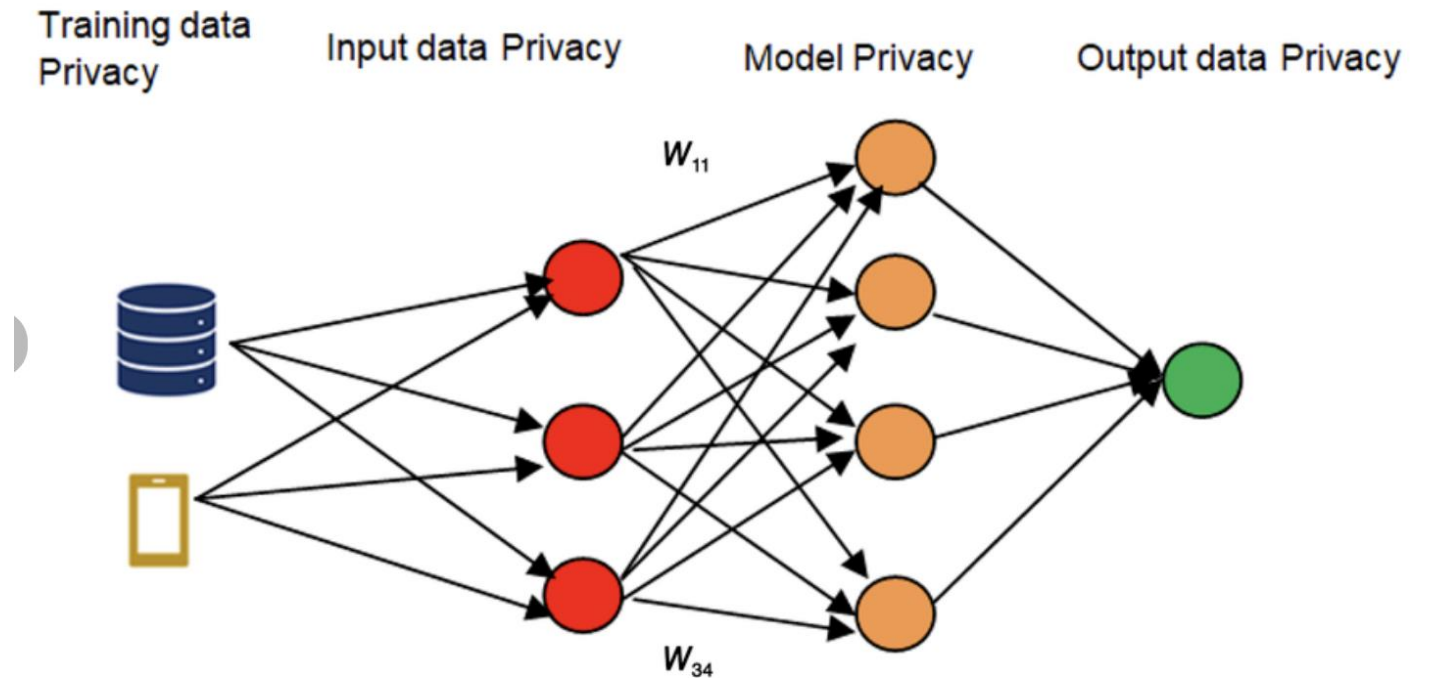
$\epsilon=1.1$ epsilon

Updated suspicion = range of
25% to 75%.

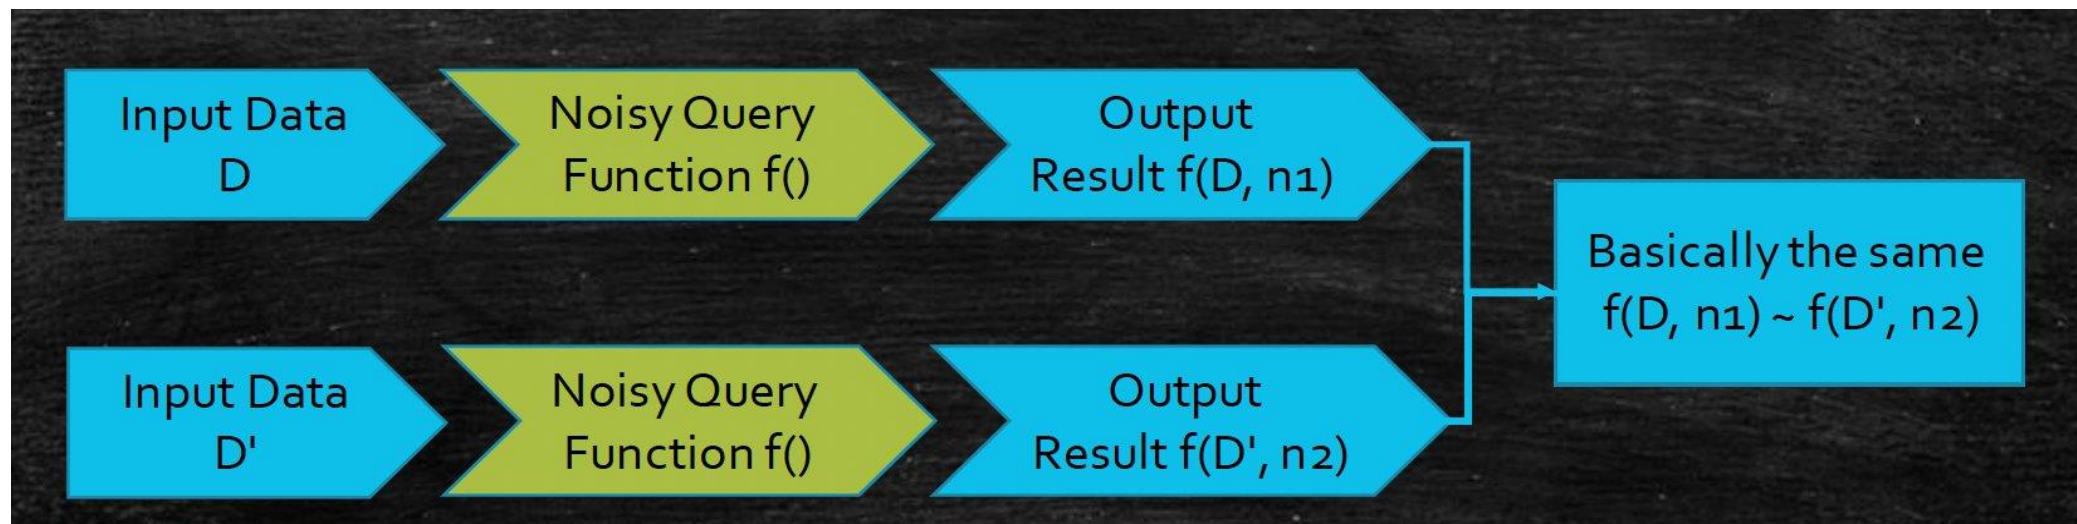
Machine Learning Pipeline




Differential Privacy Pipeline






Data privacy at different stages in deep learning network



Differential Privacy in Practice



Government, Healthcare and Technology



	Local Differential Privacy (LDP)	Central Differential Privacy (CDP)	Differentially Private SGD (DP-SGD)
Definition	Each user's data is randomized before it is sent to the server. The server never sees the raw data.	Raw data is collected and stored by a trusted server, and noise is added when releasing aggregate results.	A training algorithm where noise is added to gradients during model training to ensure privacy.
Noise	At the client/user side (before data collection).	At the server/output side (after data collection).	During model training, to the gradients and clipped updates.
Privacy Parameter (ϵ)	Typically larger ϵ (less accuracy) because each user adds noise individually.	Smaller ϵ possible with the same privacy level — better accuracy since aggregation reduces noise.	ϵ controlled via gradient clipping and noise multiplier; balance between model utility and privacy.

	Local Differential Privacy (LDP)	Central Differential Privacy (CDP)	Differentially Private SGD (DP-SGD)
Key Trade-off	Strongest privacy, weakest accuracy	Balanced privacy and accuracy	Moderate privacy, good accuracy for ML
Example Use Cases	<ul style="list-style-type: none">- Google Chrome’s RAPPOR for user telemetry- Apple’s iOS keyboard data- Privacy-preserving analytics without raw data sharing	<ul style="list-style-type: none">- Census data release (e.g., US Census Bureau)- Aggregate statistics and queries on sensitive datasets	<ul style="list-style-type: none">- Training privacy-preserving ML models- Used in Google’s and OpenAI’s research for differentially private training
Frameworks / Tools	OpenDP / Google’s RAPPOR / Apple LDP framework	TensorFlow Privacy (for central DP queries) / OpenDP	TensorFlow Privacy, PyTorch Opacus, Diffprivlib

IBM/differential-privacy-library



Diffprivlib: The IBM Differential Privacy Library

11 Contributors
279 Used by
890 Stars
207 Forks



google/rappor

RAPPOR: Privacy-Preserving Reporting Algorithms



9 Contributors
26 Issues
866 Stars
163 Forks



meta-pytorch/opacus



Training PyTorch models with differential privacy

64 Contributors
1k Used by
2k Stars
382 Forks



Introducing Tensor
Flow Privacy

OpenMined/PyDP



The Python Differential Privacy Library. Built on top of: <https://github.com/google/differential-privacy>

54 Contributors
13 Used by
541 Stars
141 Forks



VaultGemma


Differential Privacy Framework and Tools

LIBRARIES FOR PRODUCTION USE

1. Google's PipelineDP:
pip install pipeline-dp
→ Production-ready, used internally at Google
2. OpenMined's PyDP:
pip install python-dp
→ Python bindings for Google's differential-privacy library
3. IBM's diffprivlib:
pip install diffprivlib
→ Scikit-learn compatible, easy to use

References

Google



<https://github.com/sharikalog7/Differential-Privacy>